

Exploring_Ebay_Car_Sales_Data

March 6, 2020

```
[1]: import pandas as pd
      #Some .csv files can't be accessed without proper encoding method.
      autos = pd.read_csv('autos.csv', encoding='Latin-1')
```

```
[2]: #To display the dataframe.
      autos
```

```
[2]:
```

	dateCrawled	name \
0	2016-03-26 17:47:46	Peugeot_807_160_NAVTECH_ON_BOARD
1	2016-04-04 13:38:56	BMW_740i_4_4_Liter_HAMANN_UMBAU_Mega_Optik
2	2016-03-26 18:57:24	Volkswagen_Golf_1.6_United
3	2016-03-12 16:58:10	Smart_smart_fortwo_coupe_softouch/F1/Klima/Pan...
4	2016-04-01 14:38:50	Ford_Focus_1_6_Benzin_TÜV_neu_ist_sehr_gepfleg...
...
49995	2016-03-27 14:38:19	Audi_Q5_3.0_TDI_qu._S_tr._Navi_Panorama_Xenon
49996	2016-03-28 10:50:25	Opel_Astra_F_Cabrio_Bertone_Edition__TÜV_neu+...
49997	2016-04-02 14:44:48	Fiat_500_C_1.2_Dualogic_Lounge
49998	2016-03-08 19:25:42	Audi_A3_2.0_TDI_Sportback_Ambition
49999	2016-03-14 00:42:12	Opel_Vectra_1.6_16V

	seller	offerType	price	abtest	vehicleType	yearOfRegistration \
0	privat	Angebot	\$5,000	control	bus	2004
1	privat	Angebot	\$8,500	control	limousine	1997
2	privat	Angebot	\$8,990	test	limousine	2009
3	privat	Angebot	\$4,350	control	kleinwagen	2007
4	privat	Angebot	\$1,350	test	kombi	2003
...
49995	privat	Angebot	\$24,900	control	limousine	2011
49996	privat	Angebot	\$1,980	control	cabrio	1996
49997	privat	Angebot	\$13,200	test	cabrio	2014
49998	privat	Angebot	\$22,900	control	kombi	2013
49999	privat	Angebot	\$1,250	control	limousine	1996

	gearbox	powerPS	model	odometer	monthOfRegistration	fuelType \
0	manuell	158	andere	150,000km	3	lpg
1	automatik	286	7er	150,000km	6	benzin
2	manuell	102	golf	70,000km	7	benzin

3	automatik	71	fortwo	70,000km	6	benzin
4	manuell	0	focus	150,000km	7	benzin
...
49995	automatik	239	q5	100,000km	1	diesel
49996	manuell	75	astra	150,000km	5	benzin
49997	automatik	69	500	5,000km	11	benzin
49998	manuell	150	a3	40,000km	11	diesel
49999	manuell	101	vectra	150,000km	1	benzin

	brand	notRepairedDamage	dateCreated	nrOfPictures	\
0	peugeot	nein	2016-03-26 00:00:00	0	
1	bmw	nein	2016-04-04 00:00:00	0	
2	volkswagen	nein	2016-03-26 00:00:00	0	
3	smart	nein	2016-03-12 00:00:00	0	
4	ford	nein	2016-04-01 00:00:00	0	
...
49995	audi	nein	2016-03-27 00:00:00	0	
49996	opel	nein	2016-03-28 00:00:00	0	
49997	fiat	nein	2016-04-02 00:00:00	0	
49998	audi	nein	2016-03-08 00:00:00	0	
49999	opel	nein	2016-03-13 00:00:00	0	

	postalCode	lastSeen
0	79588	2016-04-06 06:45:54
1	71034	2016-04-06 14:45:08
2	35394	2016-04-06 20:15:37
3	33729	2016-03-15 03:16:28
4	39218	2016-04-01 14:38:50
...
49995	82131	2016-04-01 13:47:40
49996	44807	2016-04-02 14:18:02
49997	73430	2016-04-04 11:47:27
49998	35683	2016-04-05 16:45:07
49999	45897	2016-04-06 21:18:48

[50000 rows x 20 columns]

```
[3]: autos.info() #To display the info of the autos dataframe.
      autos.head() #Displaying top 5 rows in the dataframe.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 20 columns):
dateCrawled      50000 non-null object
name             50000 non-null object
seller          50000 non-null object
offerType       50000 non-null object
```

```

price          50000 non-null object
abtest         50000 non-null object
vehicleType    44905 non-null object
yearOfRegistration 50000 non-null int64
gearbox        47320 non-null object
powerPS        50000 non-null int64
model          47242 non-null object
odometer       50000 non-null object
monthOfRegistration 50000 non-null int64
fuelType       45518 non-null object
brand          50000 non-null object
notRepairedDamage 40171 non-null object
dateCreated    50000 non-null object
nrOfPictures   50000 non-null int64
postalCode     50000 non-null int64
lastSeen       50000 non-null object
dtypes: int64(5), object(15)
memory usage: 7.6+ MB

```

```

[3]:          dateCrawled                                     name \
0  2016-03-26 17:47:46                                Peugeot_807_160_NAVTECH_ON_BOARD
1  2016-04-04 13:38:56                        BMW_740i_4_4_Liter_HAMANN_UMBAU_Mega_Optik
2  2016-03-26 18:57:24                        Volkswagen_Golf_1.6_United
3  2016-03-12 16:58:10  Smart_smart_fortwo_coupe_softtouch/F1/Klima/Pan...
4  2016-04-01 14:38:50  Ford_Focus_1_6_Benzin_TÜV_neu_ist_sehr_gepfleg...

```

```

      seller offerType  price  abtest vehicleType  yearOfRegistration \
0  privat  Angebot  $5,000  control         bus             2004
1  privat  Angebot  $8,500  control  limousine             1997
2  privat  Angebot  $8,990    test  limousine             2009
3  privat  Angebot  $4,350  control  kleinwagen             2007
4  privat  Angebot  $1,350    test         kombi             2003

```

```

      gearbox  powerPS  model  odometer  monthOfRegistration  fuelType \
0  manuell    158  andere  150,000km             3      lpg
1  automatik   286    7er  150,000km             6  benzin
2  manuell    102  golf   70,000km             7  benzin
3  automatik    71  fortwo  70,000km             6  benzin
4  manuell     0  focus   150,000km             7  benzin

```

```

      brand  notRepairedDamage  dateCreated  nrOfPictures \
0  peugeot                nein  2016-03-26 00:00:00        0
1    bmw                nein  2016-04-04 00:00:00        0
2  volkswagen            nein  2016-03-26 00:00:00        0
3    smart                nein  2016-03-12 00:00:00        0
4    ford                nein  2016-04-01 00:00:00        0

```

	postalCode		lastSeen
0	79588	2016-04-06	06:45:54
1	71034	2016-04-06	14:45:08
2	35394	2016-04-06	20:15:37
3	33729	2016-03-15	03:16:28
4	39218	2016-04-01	14:38:50

PART ONE: Analysis on autos dataframe

- Upon looking at the information of the entire data frame, there are 5 out of a total 20 columns that have null or NaN data: `vehicle`, `gearbox`, `model`, `fuelType`, and `notRepairedDamage`. But none of the columns have more than 20% of its data as null values.
 - `vehicleType`: 5,095 null/NaN data
 - `gearbox`: 2,680 null/NaN data
 - `model`: 2,758 null/NaN data
 - `fuelType`: 4,482 null/NaN data
 - `notRepairedDamage`: 9,829 null/NaN data
- There are also 5 columns that have data that are `int` datatype: `yearOfRegistration`, `powerPS`, `monthOfRegistration`, `nrOfPictures`, `postalCode`
- Also noticed some data columns have data that's spelled incorrectly, contain a combination of numbers and characters, and datetimes.
- I expected that we will be filling null cells with appropriate data, correcting misspelled words, and removing characters from number+char combos, creating datetime objects and extracting either the dates or the times, and converting the values (which would then be numbers) to either the `int` or `float` datatypes among other tasks.
- Some cells that have strings are also written in German, so I suspect that we would have to translate all of our non-english cells to english

```
[4]: autos.columns #To display all unique column titles.
```

```
[4]: Index(['dateCrawled', 'name', 'seller', 'offerType', 'price', 'abtest',
          'vehicleType', 'yearOfRegistration', 'gearbox', 'powerPS', 'model',
          'odometer', 'monthOfRegistration', 'fuelType', 'brand',
          'notRepairedDamage', 'dateCreated', 'nrOfPictures', 'postalCode',
          'lastSeen'],
          dtype='object')
```

```
[5]: #Replacing column names to follow Python preferred snakecase format.
autos.rename(columns = {'yearOfRegistration' : 'registration_year',
                        'monthOfRegistration' : 'registration_month',
                        'notRepairedDamage' : 'unrepaired_damage',
                        'dateCreated' : 'ad_created',
                        'offerType' : 'offer_type',
                        'powerPS' : 'power_ps',
                        'nrOfPictures' : 'nr_of_pictures',
                        'postalCode' : 'postal_code',
                        'lastSeen' : 'last_seen',
                        'dateCrawled' : 'date_crawled',
```

```

        'fuelType' : 'fuel_type',
        'vehicleType' : 'vehicle_type',
        'abtest' : 'ab_test'
    }, inplace=True)
autos.columns

```

```

[5]: Index(['date_crawled', 'name', 'seller', 'offer_type', 'price', 'ab_test',
          'vehicle_type', 'registration_year', 'gearbox', 'power_ps', 'model',
          'odometer', 'registration_month', 'fuel_type', 'brand',
          'unrepaired_damage', 'ad_created', 'nr_of_pictures', 'postal_code',
          'last_seen'],
          dtype='object')

```

PART TWO: Renaming columns. - We renamed the columns whose names were in the camelcase format to the snakecase format because snakecase is the format most preferred in Python.

```

[6]: autos.describe() #Describing general information about columns in the dataframe
      ↪ containing numeric values.

```

```

[6]:      registration_year      power_ps  registration_month  nr_of_pictures  \
count      50000.000000  50000.000000      50000.000000      50000.0
mean        2005.073280    116.355920         5.723360         0.0
std          105.712813    209.216627         3.711984         0.0
min           1000.000000     0.000000         0.000000         0.0
25%           1999.000000     70.000000         3.000000         0.0
50%           2003.000000    105.000000         6.000000         0.0
75%           2008.000000    150.000000         9.000000         0.0
max           9999.000000   17700.000000        12.000000         0.0

      postal_code
count  50000.000000
mean   50813.627300
std    25779.747957
min     1067.000000
25%    30451.000000
50%    49577.000000
75%    71540.000000
max    99998.000000

```

PART THREE: Noting descriptive info. about table data - **nr_of-pictures** column has a mean of 0.0 (and everywhere else on the stats sheet) and a count of 50,000 which means that all the data in that column are 0s - **power_ps** which is another name for Horse Power(HP), has a min of 0.00, which means that there are cars registered that can't run at all. - **postal_code** column has a min value of 1067. Postal codes are meant to be 5 digits long which indicates that there are some postal codes that start with a zero which doesn't get shown. In theory, the data type in this column have to be changed to the **string** type instead of **int** or **float** types in order to show that zero in the beginning of each postal code.

```
[7]: #Removing non-numeric values from columns we want to have as numeric, and then
      ↪converting them to numeric datatypes.
autos['price'] = autos['price'].str.replace('$', '').str.replace(',', '').
      ↪astype(float)
autos['odometer'] = autos['odometer'].str.replace('km', '').str.replace(',', '').
      ↪astype(float)
autos.rename(columns={'odometer':'odometer_km'}, inplace=True)
autos.head()
```

```
[7]:
```

	date_crawled	name \
0	2016-03-26 17:47:46	Peugeot_807_160_NAVTECH_ON_BOARD
1	2016-04-04 13:38:56	BMW_740i_4_4_Liter_HAMANN_UMBAU_Mega_Optik
2	2016-03-26 18:57:24	Volkswagen_Golf_1.6_United
3	2016-03-12 16:58:10	Smart_smart_fortwo_coupe_softouch/F1/Klima/Pan...
4	2016-04-01 14:38:50	Ford_Focus_1_6_Benzin_TÜV_neu_ist_sehr_gepfleg...

	seller	offer_type	price	ab_test	vehicle_type	registration_year \
0	privat	Angebot	5000.0	control	bus	2004
1	privat	Angebot	8500.0	control	limousine	1997
2	privat	Angebot	8990.0	test	limousine	2009
3	privat	Angebot	4350.0	control	kleinwagen	2007
4	privat	Angebot	1350.0	test	kombi	2003

	gearbox	power_ps	model	odometer_km	registration_month	fuel_type \
0	manuell	158	andere	150000.0	3	lpg
1	automatik	286	7er	150000.0	6	benzin
2	manuell	102	golf	70000.0	7	benzin
3	automatik	71	fortwo	70000.0	6	benzin
4	manuell	0	focus	150000.0	7	benzin

	brand	unrepaired_damage	ad_created	nr_of_pictures \
0	peugeot	nein	2016-03-26 00:00:00	0
1	bmw	nein	2016-04-04 00:00:00	0
2	volkswagen	nein	2016-03-26 00:00:00	0
3	smart	nein	2016-03-12 00:00:00	0
4	ford	nein	2016-04-01 00:00:00	0

	postal_code	last_seen
0	79588	2016-04-06 06:45:54
1	71034	2016-04-06 14:45:08
2	35394	2016-04-06 20:15:37
3	33729	2016-03-15 03:16:28
4	39218	2016-04-01 14:38:50

```
[8]: autos['price'].unique().shape #To get the number of unique prices in the price
      ↪column.
```

```
[8]: (2357,)
```

```
[9]: autos['price'].describe() #General info about price column after numeric data
    ↪ conversion.
```

```
[9]: count      5.000000e+04
    mean      9.840044e+03
    std       4.811044e+05
    min       0.000000e+00
    25%       1.100000e+03
    50%       2.950000e+03
    75%       7.200000e+03
    max       1.000000e+08
    Name: price, dtype: float64
```

PART FOUR: Identifying and removing outliers - Based on the descriptions given above, you can already see a few outliers. For instance: the cheapest car you can get which is the `min` price value is 0. - So essentially you would be getting those cars for free. Thus, you would want to remove free cars from your dataframe because the car very low to no value. - You can also see that the `max` is another outlier, as the most expensive car in the dataframe is listed at 100 million dollars. The most expensive car in the world today costs 13 million dollars. Therefore, you would want to remove any cars priced at more than 13 million. - Ideally, the most accurate listings can be found between 1,100 and 13 million dollars in my opinion.

```
[10]: autos['price'].value_counts().head() #Displaying top 5 counts of specific
    ↪ prices from the price column.
```

```
[10]: 0.0      1421
    500.0    781
    1500.0   734
    2500.0   643
    1200.0   639
    Name: price, dtype: int64
```

```
[11]: #To find the cars with the highest prices in autos to remove.
    autos['price'].value_counts().sort_index(ascending=False).head()
```

```
[11]: 99999999.0    1
    27322222.0    1
    12345678.0    3
    11111111.0    2
    10000000.0    1
    Name: price, dtype: int64
```

```
[39]: #To find the cars with lowest prices that we don't since these cars are up for
    ↪ bids.
    autos['price'].value_counts().sort_index(ascending=True).head()
```

```
[39]: 1.0    150
      2.0     2
      3.0     1
      5.0     2
      8.0     1
      Name: price, dtype: int64
```

```
[13]: #To find unique odometer values.
      autos['odometer_km'].unique().shape
```

```
[13]: (13,)
```

```
[14]: #Describing general statistics info. on 'odometer_km' column.
      autos['odometer_km'].describe()
```

```
[14]: count      50000.000000
      mean      125732.700000
      std       40042.211706
      min       5000.000000
      25%      125000.000000
      50%      150000.000000
      75%      150000.000000
      max      150000.000000
      Name: odometer_km, dtype: float64
```

```
[15]: #Showing count of top 5 unique 'odometer_km' values in descending order
      autos['odometer_km'].value_counts().head()
```

```
[15]: 150000.0    32424
      125000.0     5170
      100000.0     2169
      90000.0      1757
      80000.0      1436
      Name: odometer_km, dtype: int64
```

```
[16]: #Showing count of top 5 unique 'odometer_km' values in descending order.
      autos['odometer_km'].value_counts().sort_index(ascending=False).head()
```

```
[16]: 150000.0    32424
      125000.0     5170
      100000.0     2169
      90000.0      1757
      80000.0      1436
      Name: odometer_km, dtype: int64
```

```
[17]: #Same data as above, just shown in ascending index order.
      autos['odometer_km'].value_counts().sort_index(ascending=True).head()
```



```
[17]: 5000.0      967
      10000.0    264
      20000.0    784
      30000.0    789
      40000.0    819
      Name: odometer_km, dtype: int64
```

```
[18]: #General statistics info on 'registration_year' column.
      autos["registration_year"].describe()
```

```
[18]: count      50000.000000
      mean        2005.073280
      std         105.712813
      min         1000.000000
      25%         1999.000000
      50%         2003.000000
      75%         2008.000000
      max         9999.000000
      Name: registration_year, dtype: float64
```

```
[19]: #Only looking at cars registered between 1900 and 2016. Anything registered
      ↪ before or after that is incorrect data.
      autos[autos["registration_year"].between(1900.0,2016.0)].describe()
```

```
[19]:
```

	price	registration_year	power_ps	odometer_km \
count	4.802800e+04	48028.00000	48028.000000	48028.000000
mean	9.585252e+03	2002.80351	117.070417	125544.161739
std	4.843817e+05	7.31085	195.151278	40106.751417
min	0.000000e+00	1910.00000	0.000000	5000.000000
25%	1.150000e+03	1999.00000	71.000000	100000.000000
50%	2.990000e+03	2003.00000	107.000000	150000.000000
75%	7.400000e+03	2008.00000	150.000000	150000.000000
max	1.000000e+08	2016.00000	17700.000000	150000.000000

	registration_month	nr_of_pictures	postal_code
count	48028.000000	48028.0	48028.000000
mean	5.767760	0.0	50935.867327
std	3.696802	0.0	25792.079828
min	0.000000	0.0	1067.000000
25%	3.000000	0.0	30459.000000
50%	6.000000	0.0	49696.000000
75%	9.000000	0.0	71665.000000
max	12.000000	0.0	99998.000000

There are a few discrepancies in the `registration_year` column. The minimum year is listed at year 1000 which is well before the first car was invented, and max year listed at year 9999 which is well into the future. Thus, we'll only looked at cars registered between 1900 - 2016 which will

remove any years less than year 1900, and years more than year 2016.

```
[20]: #Looking for percentages of cars registered based on years.
autos = autos[autos["registration_year"].between(1900.0,2016.0)]
autos["registration_year"].value_counts(normalize=True).head(10).sort_values
```

```
[20]: <bound method Series.sort_values of 2000    0.069834
2005    0.062776
1999    0.062464
2004    0.056988
2003    0.056779
2006    0.056384
2001    0.056280
2002    0.052740
1998    0.051074
2007    0.047972
Name: registration_year, dtype: float64>
```

We can see above that the majority of cars were sold between 1998 and 2016.

When we look at high price ranges for the column, we can see a significant jump from 350K dollars and up. Therefore it is safe to remove any data with prices 350K dollars or more. You can also that there are some cars less that 100 dollars. It should be safe to keep the prices of any car a dollar and up since Ebay is a site well-known for auctioning off its products.

```
[21]: #Only looking at cars between $1 and $350,000 dollars in order to narrow down
      ↳common pricing a bit better.
autos = autos[autos['price'].between(1, 350000)]
autos.shape
```

```
[21]: (46681, 20)
```

```
[22]: #General statistical info. on 'price' column.
autos['price'].describe()
```

```
[22]: count    46681.000000
mean      5977.716801
std       9177.909479
min        1.000000
25%       1250.000000
50%       3100.000000
75%       7500.000000
max      350000.000000
Name: price, dtype: float64
```

A few of our columns represent dates in the form of strings: - date_crawled - ad_created - last_seen
Here's a look at a few of rows of these columns down below

```
[23]: #Looking at date columns.  
autos[['date_crawled', 'ad_created', 'last_seen']].head()
```

```
[23]:
```

	date_crawled	ad_created	last_seen
0	2016-03-26 17:47:46	2016-03-26 00:00:00	2016-04-06 06:45:54
1	2016-04-04 13:38:56	2016-04-04 00:00:00	2016-04-06 14:45:08
2	2016-03-26 18:57:24	2016-03-26 00:00:00	2016-04-06 20:15:37
3	2016-03-12 16:58:10	2016-03-12 00:00:00	2016-03-15 03:16:28
4	2016-04-01 14:38:50	2016-04-01 00:00:00	2016-04-01 14:38:50

We're only interested in the dates, not the times.

```
[24]: #Not worried about times, the dates so we'll be taking the 1st 10 characters in  
      ↪ each of the dates column.  
(autos['date_crawled'].str[:10].value_counts(normalize=True, dropna=False).  
      ↪ sort_index())
```

```
[24]:
```

2016-03-05	0.025192
2016-03-06	0.014160
2016-03-07	0.036246
2016-03-08	0.033547
2016-03-09	0.033247
2016-03-10	0.032240
2016-03-11	0.032454
2016-03-12	0.036824
2016-03-13	0.015874
2016-03-14	0.036332
2016-03-15	0.034361
2016-03-16	0.029498
2016-03-17	0.031790
2016-03-18	0.012810
2016-03-19	0.034661
2016-03-20	0.038024
2016-03-21	0.037317
2016-03-22	0.032840
2016-03-23	0.032197
2016-03-24	0.029477
2016-03-25	0.031512
2016-03-26	0.032069
2016-03-27	0.030783
2016-03-28	0.034597
2016-03-29	0.034104
2016-03-30	0.033804
2016-03-31	0.031790
2016-04-01	0.033804
2016-04-02	0.035561
2016-04-03	0.038774

```

2016-04-04    0.036610
2016-04-05    0.013003
2016-04-06    0.003085
2016-04-07    0.001414
Name: date_crawled, dtype: float64

```

Let's do the same for our `ad_created` column

```

[25]: #Looking at percentage of cars registered based on unique dates in 'ad_created'
      ↪column.
      (autos['ad_created'].str[:10].value_counts(normalize=True, dropna=False).
      ↪sort_index())

```

```

[25]: 2015-06-11    0.000021
      2015-08-10    0.000021
      2015-09-09    0.000021
      2015-11-10    0.000021
      2015-12-05    0.000021
      ...
      2016-04-03    0.039009
      2016-04-04    0.036953
      2016-04-05    0.011782
      2016-04-06    0.003170
      2016-04-07    0.001264
      Name: ad_created, Length: 74, dtype: float64

```

Now for our `last_seen` column

```

[26]: #Looking at percentage of cars registered based on unique dates in 'ad_created'
      ↪column.
      (autos['last_seen'].str[:10].value_counts(normalize=True, dropna=False).
      ↪sort_index())

```

```

[26]: 2016-03-05    0.001071
      2016-03-06    0.004113
      2016-03-07    0.005377
      2016-03-08    0.007476
      2016-03-09    0.009768
      2016-03-10    0.010690
      2016-03-11    0.012382
      2016-03-12    0.023757
      2016-03-13    0.008654
      2016-03-14    0.012660
      2016-03-15    0.016002
      2016-03-16    0.016281
      2016-03-17    0.028084
      2016-03-18    0.007219
      2016-03-19    0.015617

```

```

2016-03-20    0.020629
2016-03-21    0.020587
2016-03-22    0.020844
2016-03-23    0.018359
2016-03-24    0.019687
2016-03-25    0.018937
2016-03-26    0.016795
2016-03-27    0.015638
2016-03-28    0.020694
2016-03-29    0.022086
2016-03-30    0.024614
2016-03-31    0.023628
2016-04-01    0.022943
2016-04-02    0.024657
2016-04-03    0.025149
2016-04-04    0.024121
2016-04-05    0.125404
2016-04-06    0.223324
2016-04-07    0.132752
Name: last_seen, dtype: float64

```

The `last_seen` date column shows a spike in the last 3 days of sales. This is most likely due to the bidding war strategy when bidders typically wait until the last few days or the last day to make their final bids. The days prior can't have any relevant effect since the percentages are pretty evenly distributed.

```

[27]: #Looking at percentage of cars registered based on unique brands.
autos['brand'].value_counts(normalize=True)

```

```

[27]: volkswagen    0.211264
      bmw          0.110045
      opel         0.107581
      mercedes_benz 0.096463
      audi         0.086566
      ford         0.069900
      renault      0.047150
      peugeot      0.029841
      fiat         0.025642
      seat         0.018273
      skoda        0.016409
      nissan        0.015274
      mazda        0.015188
      smart        0.014160
      citroen      0.014010
      toyota       0.012703
      hyundai      0.010025
      sonstige_autos 0.009811

```

volvo	0.009147
mini	0.008762
mitsubishi	0.008226
honda	0.007840
kia	0.007069
alfa_romeo	0.006641
porsche	0.006127
suzuki	0.005934
chevrolet	0.005698
chrysler	0.003513
dacia	0.002635
daihatsu	0.002506
jeep	0.002271
subaru	0.002142
land_rover	0.002099
saab	0.001649
jaguar	0.001564
daewoo	0.001500
trabant	0.001392
rover	0.001328
lancia	0.001071
lada	0.000578

Name: brand, dtype: float64

The top 5 car brands on this list are all German made. The top German brand more than doubles the next car brand from the next country. We'll limit our analysis to brands that accounts for more than 5% of the total sales data

```
[28]: brands = autos['brand'].value_counts(normalize = True)
      most_common_brands = brands[brands > .05].index
      most_common_brands
```

```
[28]: Index(['volkswagen', 'bmw', 'opel', 'mercedes_benz', 'audi', 'ford'],
      dtype='object')
```

```
[29]: common_brand_dict={}
      for brand in most_common_brands:
          selected_rows = autos[autos["brand"]==brand]
          mean_price = selected_rows["price"].mean()
          common_brand_dict[brand] = round(mean_price, 2)
      import operator
      brand_dict_sorted = sorted(common_brand_dict.items(), key=operator.
      ↪itemgetter(1), reverse=True)
      print('Average price for each car brand in the top 6 in descending order:',
      ↪brand_dict_sorted)
```

```
Average price for each car brand in the top 6 in descending order: [('audi',
9336.69), ('mercedes_benz', 8628.45), ('bmw', 8332.82), ('volkswagen', 5402.41),
```

```
('ford', 3749.47), ('opel', 2975.24)]
```

As we can see, the cheapest commonly sold brands are **ford** and **opel**.

The most expensive commonly sold brands are **audi** and **mercedes_benz**.

The car brands commonly sold that are priced in between are **bmw** and **volkswagen**.

Out of the top 6 cars on the list, volkswagens are the most commonly sold car although it is priced in between which shows that customers not only value saving money, but they also value quality as well. Volkswagen cars are known for their top quality, safety, and engineering.

```
[30]: #Create a series for common_brand_dict.  
      bmp_series = pd.Series(common_brand_dict)  
      print(bmp_series)
```

```
volkswagen    5402.41  
bmw           8332.82  
opel          2975.24  
mercedes_benz 8628.45  
audi          9336.69  
ford          3749.47  
dtype: float64
```

```
[31]: #Convert common_brand_dict series to a Dataframe.  
      mean_price_df = pd.DataFrame(bmp_series, columns=['mean_price'])  
      mean_price_df
```

```
[31]:          mean_price  
volkswagen    5402.41  
bmw           8332.82  
opel          2975.24  
mercedes_benz 8628.45  
audi          9336.69  
ford          3749.47
```

```
[32]: avg_mileage_dict = {}  
      for brand in most_common_brands:  
          selected_row = autos[autos['brand'] == brand]  
          #Converting Kilometers to Miles.  
          mileage = (selected_row['odometer_km'] / 1.609)  
          mean_mileage = mileage.mean()  
          #Round values to 2 decimal places.  
          avg_mileage_dict[brand] = round(mean_mileage, 2)
```

```
[33]: #Create a series for avg_mileage_dict  
      avg_mileage_series = pd.Series(avg_mileage_dict)  
      avg_mileage_series
```

```
[33]: volkswagen      79992.02
      bmw            82394.35
      opel           80366.71
      mercedes_benz  81285.50
      audi           80271.84
      ford           77231.83
      dtype: float64
```

```
[34]: #Convert avg_mileage_dict series to a Dataframe.
avg_mileage_df = pd.DataFrame(avg_mileage_series, columns=['avg_mileage'])
avg_mileage_df
```

```
[34]:          avg_mileage
volkswagen      79992.02
bmw             82394.35
opel            80366.71
mercedes_benz   81285.50
audi            80271.84
ford            77231.83
```

```
[35]: #Combine both Dataframes
pd.concat([mean_price_df, avg_mileage_df], axis=1)
```

```
[35]:          mean_price  avg_mileage
volkswagen      5402.41      79992.02
bmw             8332.82      82394.35
opel            2975.24      80366.71
mercedes_benz   8628.45      81285.50
audi            9336.69      80271.84
ford            3749.47      77231.83
```

We observe that `avg_mileage` doesn't vary as much as the `avg_price`. We can see that the more expensive brands (BMW, Audi, and Mercedes Benz) tend to have higher mileages on average than the cheaper brands with Opel being the only exception.

```
[36]: #Replacing German words with their English equivalents.
autos['seller'] = autos['seller'].replace('privat', 'private')
autos['offer_type'] = autos['offer_type'].replace('Angebot', 'Offer')
autos['vehicle_type'] = autos['vehicle_type'].replace('kleinwagen', 'small_car')
autos['vehicle_type'] = autos['vehicle_type'].replace('kombi', 'station_wagon')
autos['vehicle_type'] = autos['vehicle_type'].replace('cabrio', 'convertible')
autos['vehicle_type'] = autos['vehicle_type'].replace('andere', 'other')
autos['gearbox'] = autos['gearbox'].replace('automatik', 'automatic')
autos['gearbox'] = autos['gearbox'].replace('manuell', 'manual')
autos['unrepaired_damage'] = autos['unrepaired_damage'].replace('nein', 'no')
autos['unrepaired_damage'] = autos['unrepaired_damage'].replace('ja', 'yes')
```

Our dates columns currently contain dates and times. We want those columns to contain dates

only and we also want to remove all the -s in those dates which will allow us to convert the dates to integers.

```
[37]: #Retrieving just the date and not the times and then converting them to type_
      ↪ integers
autos['last_seen'] = autos['last_seen'].str[:10]
autos['date_crawled'] = autos['date_crawled'].str[:10]
autos['ad_created'] = autos['ad_created'].str[:10]
autos['last_seen'] = autos['last_seen'].str.replace('-', '').astype(int)
autos['date_crawled'] = autos['date_crawled'].str.replace('-', '').astype(int)
autos['ad_created'] = autos['ad_created'].str.replace('-', '').astype(int)
autos
```

```
[37]:
```

	date_crawled	name \
0	20160326	Peugeot_807_160_NAVTECH_ON_BOARD
1	20160404	BMW_740i_4_4_Liter_HAMANN_UMBAU_Mega_Optik
2	20160326	Volkswagen_Golf_1.6_United
3	20160312	Smart_smart_fortwo_coupe_softouch/F1/Klima/Pan...
4	20160401	Ford_Focus_1_6_Benzin_TÜV_neu_ist_sehr_gepfleg...
...
49995	20160327	Audi_Q5_3.0_TDI_qu._S_tr._Navi_Panorama_Xenon
49996	20160328	Opel_Astra_F_Cabrio_Bertone_Edition__TÜV_neu+...
49997	20160402	Fiat_500_C_1.2_Dualogic_Lounge
49998	20160308	Audi_A3_2.0_TDI_Sportback_Ambition
49999	20160314	Opel_Vectra_1.6_16V

	seller	offer_type	price	ab_test	vehicle_type	registration_year \
0	private	Offer	5000.0	control	bus	2004
1	private	Offer	8500.0	control	limousine	1997
2	private	Offer	8990.0	test	limousine	2009
3	private	Offer	4350.0	control	small_car	2007
4	private	Offer	1350.0	test	station_wagon	2003
...
49995	private	Offer	24900.0	control	limousine	2011
49996	private	Offer	1980.0	control	convertible	1996
49997	private	Offer	13200.0	test	convertible	2014
49998	private	Offer	22900.0	control	station_wagon	2013
49999	private	Offer	1250.0	control	limousine	1996

	gearbox	power_ps	model	odometer_km	registration_month	fuel_type \
0	manual	158	andere	150000.0	3	lpg
1	automatic	286	7er	150000.0	6	benzin
2	manual	102	golf	70000.0	7	benzin
3	automatic	71	fortwo	70000.0	6	benzin
4	manual	0	focus	150000.0	7	benzin
...
49995	automatic	239	q5	100000.0	1	diesel

49996	manual	75	astra	150000.0	5	benzin
49997	automatic	69	500	5000.0	11	benzin
49998	manual	150	a3	40000.0	11	diesel
49999	manual	101	vectra	150000.0	1	benzin

	brand	unrepaired_damage	ad_created	nr_of_pictures	postal_code	\
0	peugeot	no	20160326	0	79588	
1	bmw	no	20160404	0	71034	
2	volkswagen	no	20160326	0	35394	
3	smart	no	20160312	0	33729	
4	ford	no	20160401	0	39218	
...	
49995	audi	no	20160327	0	82131	
49996	opel	no	20160328	0	44807	
49997	fiat	no	20160402	0	73430	
49998	audi	no	20160308	0	35683	
49999	opel	no	20160313	0	45897	

	last_seen
0	20160406
1	20160406
2	20160406
3	20160315
4	20160401
...	...
49995	20160401
49996	20160402
49997	20160404
49998	20160405
49999	20160406

[46681 rows x 20 columns]

Now let's take a look at whether or not there are price discrepancies based on whether or not cars with histories of damages have been repaired. We'll do this using a dictionary.

```
[38]: yes_no = ['yes', 'no']
      #Creating a dictionary for damaged prices
      damaged_prices = {}
      #Finding the average price of each car based on whether or not they are damaged,
      ↪and then storing them in dictionary
      #before making final analysis.
      for answer in yes_no:
          selected_rows = autos[autos['unrepaired_damage'] == answer]
          mean_price = selected_rows['price'].mean()
          damaged_prices[answer] = round(mean_price, 2)
      damaged_prices
```

```
[38]: {'yes': 2241.15, 'no': 7164.03}
```

We can see that cars with unrepaired damage are much cheaper than cars without damage on average.